


How I Approach Markets

START WITH THE QUESTION, NOT THE TRADE

Most market stories are easy to explain after the fact. The harder part is deciding which new information actually changes the odds. My process is to start with the question, work out what the market already expects, trace the mechanism into cash flows or prices, and then let the evidence decide whether the original idea survives.



The working rule: "this matters" and "this creates a trade" are not the same conclusion. A useful market view needs a mechanism, evidence that can change the conclusion, and a clear sense of what the original thesis missed.

Three investigations, one approach

CORPORATE FINANCING

ReNew Energy

Higher yields raised the question of whether financing costs could begin to compete with operating growth. I followed the debt structure, hedges and disclosed rate sensitivity to see how much of that risk actually remained.

AFRICAN FX RESEARCH

Forecast risk under regime change

A simple hypothesis about stale currencies failed. That pushed the research toward a harder question: when the exchange-rate regime changes, can the model's own confidence still be trusted?

MACRO / EVENT RESEARCH

Rhine low water

A physical bottleneck looked like an obvious short thesis. The operating evidence was real, but company outcomes showed that logistics redundancy mattered more than headline exposure.

WHAT I VALUE IN RESEARCH

Mechanism over narrative. I want to know how a headline reaches a balance sheet, a cash flow or a price.

Evidence over attachment. If the obvious trade does not survive the data, that is still a useful result.

ABOUT THIS SAMPLE

These are compact versions of three independent 2026 investigations. Together they show how I move from a live market question to company disclosures, market data and operating evidence - and how I revise the conclusion when the evidence points somewhere else.

[Supporting notes and research materials on GitHub](#)



ReNew Energy: When the Cost of Money Enters the Growth Story

Rising bond yields made financing-sensitive businesses worth another look. ReNew was a useful case because rapid, capital-intensive growth sits alongside floating-rate debt, foreign-currency funding and a visible refinancing schedule.

WHY I STARTED HERE

Renewable developers have an unusually direct link between growth and financing: more capacity generally means more project capex, so the cost of money can change the economics of otherwise attractive growth.

ReNew stood out because its disclosures were detailed enough to test the question properly rather than infer it from a headline debt number.

The question:

Was ReNew becoming meaningfully exposed to higher benchmark rates, or did the combination of fixed debt and hedging keep basis-point risk contained enough that the operating story still mattered more?

What I investigated before forming a view

Fixed vs. variable debt



INR vs. USD funding




SOFR / TONA exposure and hedges




Maturities and disclosed +50bp sensitivity

Figure 1 - Debt mix




Source: ReNew disclosures reviewed Sep 2026.

Figure 2 - Funding currency



Source: ReNew disclosures reviewed Sep 2026.

Figure 3 - Near-term amortisation



Total: INR 153.6bn

Source: ReNew maturity disclosures; totals may reflect rounding.

What the disclosures showed

ReNew was not an unhedged floating-rate bet. Fixed-rate debt and a broad hedge book reduced the immediate exposure, while the remaining variable debt, benchmark-linked borrowing and refinancing schedule preserved a meaningful - but more delayed - sensitivity to higher-for-longer rates.

That changed the framing from **"higher yields will squeeze ReNew now"** to **"persistent rates matter through residual floating exposure and future refinancing."**

Disclosed sensitivity: a +50bp move in rates corresponded to roughly **INR 1.98bn lower pre-tax earnings**. I treated that as evidence of residual rate sensitivity, not as a forecast of realised losses.

Conclusion: the original thesis was tempered rather than rejected. ReNew remained rate-sensitive, but the evidence did not support treating it as an acute basis-point-risk story. The more interesting risk was whether persistently high rates could gradually raise interest costs, refinancing costs and project hurdle rates enough to make future growth less accretive.

Source basis: ReNew Energy Global company disclosures and financial statements reviewed in September 2026; calculations and interpretation are independent research.

AfriFin: The First Hypothesis Failed. That Made the Research Better.

I started with a simple idea about stale African currencies and zero-shot forecasting. The data did not support it. The more interesting question turned out to be what happens to forecast confidence when the exchange-rate regime itself changes.

WHAT BIRTHED THE RESEARCH

Zero-shot time-series foundation models were being promoted as broadly generalisable forecasters. African FX markets looked like a useful stress test: some series spend long periods barely moving, then reprice sharply when policy or exchange-rate regimes change. My initial hypothesis was that this price staleness would make those currencies especially difficult for a zero-shot model.

34

African USD exchange-rate series in the accepted study

3

Foundation-model families in the post-acceptance extension

15-22pp

Approx. shortfall in nominal 80% coverage after documented regime changes

STAGE 1 - THE ACCEPTED PAPER

Evaluating Zero-Shot Time-Series Foundation Models in African Currency Markets

Accepted at GlobalSouthAI @ NeurIPS 2026 (Affinity Event).

- Rolling-origin evaluation across 34 African USD exchange-rate series using Chronos-Bolt and persistence baselines.
- The original staleness hypothesis did **not** survive: aggregate currency-level staleness did not reliably identify relative forecast failure.
- Near-constant series exposed another problem - relative-error metrics could become misleading when the denominator was close to zero.

STAGE 2 - WHAT I CHANGED AFTER ACCEPTANCE

From "which currencies fail?" to "when should confidence be trusted?"

I rebuilt the question rather than rescue the original hypothesis.

- Expanded to Chronos-Bolt, TimesFM 2.5 and Moirai 2.0.
- Added a frozen IMF AREAER regime layer, origin-level diagnostics and probabilistic calibration.
- Pre-declared the model stack, forecast origins and inferential plan before inspecting the new results.

The stronger result: after documented exchange-regime reclassifications, nominal 80% predictive intervals became sharply under-covered across all three model families. The intervals did not widen enough to reflect the structural change.

How the research evolved



What I took from it: the interesting question was no longer which model looked best on average. It was when a model's reported uncertainty stopped being informative. That is a more useful question for markets, where structural breaks are exactly when false confidence is most expensive.

Source basis: accepted AfriFin short paper and the separately versioned post-acceptance extension. The extension is not presented as part of the accepted submission.

Rhine Low Water: From an Obvious Short to a Better State Variable

The investigation started because the physical stress looked hard to ignore: water levels collapsed, freight costs exploded and industrial disruptions began to appear. The question was whether that shock was clean enough to translate into a listed-equity trade.


Original August hypothesis:

low water → lower barge capacity → road/rail substitution → spare-capacity saturation → inventory drawdown → production and margin pressure.

WHY IT WAS WORTH INVESTIGATING

Kaub fell from 150cm in early June to 25cm by 31 July and reached single digits in August. At the same time, same-route ARA-Basel freight moved from EUR35/t in early June to EUR276.67/t by mid-August. That was large enough to test whether the bottleneck had crossed from a logistics story into a balance-sheet variable.

Figure 4 - Selected Kaub gauge observations (cm)



Red zone marks readings below the 77cm GLW reference. *ELWIS intraday print; gauge level is not navigable depth.

Figure 5 - ARA-Basel freight snapshots (EUR/t)



Same-route public market snapshots. September remained far above early-summer levels after the August spike.

Did the original mechanism show up?

- 1. OPERATIONS** Yes. **Covestro declared force majeure, Evonik reported disruption at Marl, thyssenkrupp reduced blast-furnace output, and Salzgitter/HKM shifted coal flows to rail.** The physical chain had clearly crossed into company operations.
- 2. COMPARISON** But the same river shock produced very different financial sensitivity. BASF had more low-water and modal redundancy; thyssenkrupp had a more concentrated earnings exposure; ArcelorMittal's Duisburg site was too small within the group for a clean listed-equity expression.
- 3. STATE VARIABLE** That changed the thesis. **Logistics redundancy - alternative transport, inventories, low-draft capacity and pricing power - explained more than the gauge reading alone.** Low water was the trigger; redundancy determined how much of the shock reached earnings.

What broke the obvious short: across five pre-specified equity windows there was no clean, monotonic negative Rhine signal. Earnings, restructuring, analyst actions and sector policy repeatedly dominated short-horizon price moves. The result was not "the Rhine did not matter"; it was that **physical exposure alone was not a reliable short signal.**

Source basis: ELWIS/BfG, Contargo, S&P Global Commodity Insights, Reuters and company disclosures as documented in the full 25 September 2026 Rhine research note.